

**MAJOR PROJECT REPORT
ON
Ledger: AI-Based Fake News Detection System
Using Explainable AI**

Submitted By

**ANANDHAKRISHNAN S (CEC23CS033)
ARAVIND A KAMATH (CEC23CS041)
ARJUN MANOJ (CEC23CS043)
SANGEETH KRISHNAN S (CEC23CS105)**

under the esteemed guidance of

Mrs. Rajeswary R
Assistant Professor
Department of Computer Science and Engineering



SEPTEMBER 2026
DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
COLLEGE OF ENGINEERING, PALLIPPURAM P O, CHERTHALA,
ALAPPUZHA PIN: 688541,
PHONE: 0478 2553416, FAX: 0478 2552714
<http://www.cectl.ac.in>

MAJOR PROJECT REPORT ON
Ledger: AI-Based Fake News Detection System
Using Explainable AI

Submitted By

ANANDHAKRISHNAN S (CEC23CS033)
ARAVIND A KAMATH (CEC23CS041)
ARJUN MANOJ (CEC23CS043)
SANGEETH KRISHNAN S (CEC23CS105)

under the esteemed guidance of

Mrs. Rajeswary R

In partial fulfillment of the requirements for the award of the degree of

Bachelor of Technology

in

Computer Science and Engineering

of

APJ Abdul Kalam Technological University



SEPTEMBER 2026

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
COLLEGE OF ENGINEERING, PALLIPPURAM P O, CHERTHALA,
ALAPPUZHA PIN: 688541,

PHONE: 0478 2553416, FAX: 0478 2552714

<http://www.cectl.ac.in>

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
COLLEGE OF ENGINEERING CHERTHALA
ALAPPUZHA-688541



C E R T I F I C A T E

This is to certify that, the project report titled **LEDGER: AI-BASED FAKE NEWS DETECTION SYSTEM USING EXPLAINABLE AI**, is a bonafide record of the CSD 415 Project Phase I presented by **ANANDHAKRISHNAN S** (Reg. No. CEC23CS033), **ARAVIND A KAMATH** (Reg. No. CEC23CS041), **ARJUN MANOJ** (Reg. No. CEC23CS043), **SANGEETH KRISHNAN S** (Reg. No. CEC23CS105), Seventh Semester B. Tech. Computer Science and Engineering students, under our guidance and supervision, in partial fulfillment of the requirements for the award of the degree, B. Tech. Computer Science and Engineering of APJ Abdul Kalam Technological University.

Guide

Mrs. Rajeswary R

Assistant Professor

Dept. of Computer Science
& Engineering

Co-ordinator

Mrs. Rakhi R

Assistant Professor

Dept. of Computer Science
& Engineering

HOD

Dr. Preetha Theresa Joy

Professor

Dept. of Computer Science
& Engineering

ACKNOWLEDGEMENT

This work would not have been possible without the support of many people. First and foremost, we give thanks to Almighty God who gave us the inner strength, resource and ability to complete our project successfully.

We would like to thank **Dr. Jaya V.L**, our Principal, who has provided us with the best facilities and atmosphere for the project completion and presentation. We would also like to thank our HOD **Dr. Preetha Theresa Joy** (Professor, Department of Computer Science and Engineering), our Project Co-ordinator **Mrs. Rakhi R** (Assistant Professor, Department of Computer Science and Engineering) and our guide **Mrs. Rajeswary R** (Assistant Professor, Department of Computer Science and Engineering) for the help extended and also for the encouragement and support given to us while doing the project.

We would like to thank our dear friends for extending their cooperation and encouragement throughout the project work, without which we would never have completed the project this well. Thank you all for your love and also for being very understanding.

ABSTRACT

The rapid proliferation of fake news and misinformation across social media and digital platforms has emerged as a critical societal challenge, distorting democratic elections, propagating harmful medical myths and inciting communal unrest. While existing automated detection systems achieve high predictive accuracy, they predominantly operate as opaque black boxes that yield classifications without explaining the underlying reasoning, thereby undermining user trust and institutional credibility. To address these critical limitations, this project proposes **Ledger**, an Explainable AI-assisted web application for real-time fake news verification. The core classification architecture integrates **FastText** subword embeddings to robustly capture semantic representations while effectively handling out-of-vocabulary terms, informal internet slang and typographical errors. A **1D Convolutional Neural Network (CNN)** extracts salient local n-gram features and deceptive syntactic patterns, which are sequentially processed by a **Long Short-Term Memory (LSTM)** network to model long-range contextual dependencies across the text. To overcome black-box opacity, the system incorporates **LIME** (Local Interpretable Model-agnostic Explanations) to generate word-level attribution heatmaps that visually explain individual veracity decisions alongside calibrated confidence scores. Grounded in the methodology of the base paper, *Advancing Fake News Detection: Hybrid Deep Learning With FastText and Explainable AI* (Hashmi et al., IEEE Access, 2024), the proposed implementation pairs a high-performance Python **FastAPI** backend for sub-second inference with an interactive **React.js** web dashboard. This Phase I report details the problem formulation, a comprehensive literature survey and the architectural system design diagrams without claiming completed deployment or measured project performance.

Contents

1	INTRODUCTION	1
2	PROBLEM STATEMENT	3
2.1	Problem Statement	3
2.2	Objectives	4
3	LITERATURE SURVEY	5
3.1	Advancing Fake News Detection: Hybrid Deep Learning With FastText and Explainable AI	5
3.2	WELFake: Word Embedding Over Linguistic Features for Fake News Detection	6
3.3	FakeNewsNet: A Data Repository with Diverse Features in Social Media for Fake News Detection	6
3.4	Impact of Convolutional Neural Network and FastText Embedding on Text Classification	7
3.5	FakeBERT: Fake News Detection in Social Media using BERT and Deep Learning	7
3.6	A Systematic Survey on Explainable AI Applied to Fake News Detection	8
3.7	Exploiting Stance Similarity and Graph Neural Networks for Fake News Detection	8
3.8	Big Data ML-Based Fake News Detection Using Distributed Learning	9
3.9	“Liar, Liar Pants on Fire”: A New Benchmark Dataset for Fake News Detection	9
3.10	Fake News Detection: A Hybrid CNN-RNN Based Approach	10
3.11	Supervised Learning for Fake News Detection: An Explainable Approach	10
4	SYSTEM DESIGN	12
4.1	System Architecture	12
4.2	Data Flow Diagram Level 0	13

4.3	Data Flow Diagram Level 1	14
4.4	UML Sequence Diagram	15
REFERENCES		17

List of Figures

4.1	System Architecture Diagram	13
4.2	Data Flow Diagram: Level 0 (Context Diagram)	14
4.3	Data Flow Diagram: Level 1	15
4.4	UML Sequence Diagram	16

Chapter 1

INTRODUCTION

Digital platforms and social media networks have become the primary sources of news for much of the world, enabling the rapid and largely unchecked spread of misinformation. Traditional fact-checking relies heavily on manual human verification, which is slow, labour-intensive and incapable of keeping pace with viral content. The consequences of unchecked fake news are serious: it distorts democratic elections, accelerates the spread of dangerous medical myths and has repeatedly incited communal violence. Both the United Nations and the World Health Organization have classified misinformation and information epidemics as critical threats to public health and democratic governance. Developing nations face acute vulnerability because large-scale internet adoption has outpaced digital media literacy.

Most contemporary automated classifiers achieve high predictive accuracy but operate as opaque black boxes that offer no insight into why a particular article was labelled real or fake. This opacity prevents journalists, fact-checkers and the general public from evaluating or trusting a verdict. The selected base paper, *Advancing Fake News Detection: Hybrid Deep Learning With FastText and Explainable AI* (Hashmi et al., IEEE Access, 2024), addresses both problems simultaneously. It combines FastText subword embeddings with a hybrid 1D-CNN and LSTM network, achieving up to 99 % accuracy on benchmark datasets and applies LIME to provide word-level attribution that makes the model reasoning visible.

The reviewed literature explores complementary methodologies. Verma et al. demonstrate the value of combining linguistic statistics with word embeddings in the WELFake corpus. Shu et al. provide a multi-dimensional social-context repository in FakeNewsNet. Umer et al. validate FastText subword representations for noisy text classification. Transformer architectures such as FakeBERT deliver strong semantic modelling but carry high computational and latency costs. Kumar and Chacko survey Explainable AI frameworks and underline the necessity of visual decision justifications. Additional works examine stance graphs, distributed learning pipelines, truthfulness benchmarks, hybrid sequential networks and supervised linguistic explainability.

The proposed project, **Ledger: AI-Based Fake News Detection System Using**

Explainable AI, combines lightweight deep learning with interactive decision explainability. Rather than relying on external social-media graph data or expensive GPU infrastructure, Ledger operates as a standalone, text-based verification tool. Users submit a headline or full article text through a web interface and receive an instant authenticity verdict accompanied by highlighted word-level explanations.

The planned system consists of a FastText embedding layer, a 1D-CNN for local feature extraction, an LSTM for sequence modelling and a LIME attribution layer. A FastAPI server handles model serving and a React.js dashboard presents the result with confidence scores and highlighted keywords. The overall goal is to provide everyday citizens, journalists and fact-checkers with an accessible, interpretable and trustworthy fake-news verification assistant.

Chapter 2

PROBLEM STATEMENT

2.1 Problem Statement

The unchecked spread of fabricated news on digital platforms harms democratic processes, public health initiatives and social harmony. Existing automated detection approaches share several critical limitations.

First, nearly all current deep-learning classifiers function as black boxes. Even when they achieve high statistical accuracy they do not explain which words or phrases drove the prediction. Without interpretable justifications, users, journalists and regulatory bodies cannot assess whether a verdict is based on genuine linguistic indicators of deception or on spurious correlations in the training data.

Second, state-of-the-art contextual language models such as BERT require large GPU memories and produce high inference latency. These requirements make them impractical for low-cost, real-time consumer web applications where response times must remain sub-second on standard server hardware.

Third, social-media text is characterized by informal language, spelling errors, evolving internet slang and deliberate adversarial obfuscation. Standard word-embedding techniques assign fixed representations to complete tokens and fail when presented with out-of-vocabulary terms.

Fourth, many academic detection systems depend on external social-context attributes such as sharing graphs, retweet trees and comment sentiment. These dependencies make the models ineffective for newly published or standalone articles that lack prior social engagement.

The project therefore proposes Ledger, an AI-assisted web platform that addresses these limitations by combining FastText subword embeddings, a hybrid CNN-LSTM deep-learning model and LIME explainability to deliver real-time, text-only news verification with clear, word-level justifications.

2.2 Objectives

1. To design and develop an accessible web-based fake news detection system for real-time verification of news articles and headlines.
2. To use FastText subword n-gram embeddings to handle out-of-vocabulary tokens, typographical errors and informal digital slang.
3. To extract localized semantic features and deceptive phrase patterns using a 1D Convolutional Neural Network.
4. To model long-range sequential contextual dependencies using a Long Short-Term Memory network.
5. To integrate Explainable AI using LIME to produce transparent, word-level visual justifications for every classification verdict.
6. To evaluate and benchmark the hybrid framework on recognized open-access datasets including WELFake, FakeNewsNet and LIAR.
7. To construct a lightweight, high-throughput REST API using Python FastAPI for model serving and sub-second inference.
8. To develop an intuitive React.js user dashboard that displays the verdict, confidence score and an interactive highlighted explanation heatmap.

Chapter 3

LITERATURE SURVEY

The literature survey examines the selected base paper and ten complementary research works, focusing on hybrid deep learning, subword representations, contextual language models, social graph analytics, distributed learning and explainable AI for fake news detection. The limitations in each study identify technical gaps that motivated the design choices adopted in the Ledger platform.

3.1 Advancing Fake News Detection: Hybrid Deep Learning With FastText and Explainable AI

Authors: Ehtesham Hashmi, Sule Yildirim Yayilgan, Muhammad Mudassar Yamin, Subhan Ali and Mohamed Abomhara

Published On: IEEE Access, Volume 12, 2024

- **About:** Investigates the integration of hybrid deep learning with subword embeddings and LIME explainability to address the black-box limitation in fake news detection.
- **Methodology:** Text is preprocessed through regex cleaning, lemmatization and tokenization. FastText subword embeddings represent the tokens. A 1D-CNN extracts local n-gram features, an LSTM captures sequential dependencies and a Softmax layer classifies the text. LIME provides post-hoc word-level attribution.
- **Advantages:** Achieves 99 % accuracy on WELFake, 97 % on FakeNewsNet and 99 % on FakeNewsPrediction. FastText handles out-of-vocabulary terms and noisy social media text. LIME provides transparent word-level explanations that build user trust.
- **Limitations:** Evaluated primarily on English corpora; lacks multilingual support. Sensitive to superficial proper nouns and date references in short statements. Does not handle multimodal content such as manipulated images or synthetic video.

3.2 WELFake: Word Embedding Over Linguistic Features for Fake News Detection

Authors: Pawan Kumar Verma, Prateek Agrawal, Ivone Amorim and Radu Prodan

Published On: IEEE Transactions on Computational Social Systems, 2021

- **About:** Introduces WELFake, a large benchmark dataset and evaluates a hybrid framework combining linguistic statistics with dense word embeddings for fake news classification.
- **Methodology:** Aggregates 72,134 news records from four repositories. Concatenates 20 linguistic attributes with GloVe and Word2Vec embeddings and feeds the combined representation into ensemble and neural classifiers.
- **Advantages:** Provides one of the largest balanced open-access benchmark datasets, reducing overfitting. Achieves 96.73 % accuracy, outperforming standalone BERT and CNN by up to 4.25 %. Demonstrates strong cross-domain generalisation.
- **Limitations:** Manual extraction of 20 linguistic features is computationally expensive. Individual predictions lack interpretability explanations. Fixed embeddings cannot represent newly coined internet terminology.

3.3 FakeNewsNet: A Data Repository with Diverse Features in Social Media for Fake News Detection

Authors: Kai Shu, Deepak Mahudeswaran, Suhang Wang, Dongwon Lee and Huan Liu

Published On: Big Data, Mary Ann Liebert, 2020

- **About:** Presents a comprehensive data repository capturing news content, social network interactions and spatiotemporal propagation dynamics for fake news research.
- **Methodology:** Compiles verified news instances from PolitiFact and GossipCop, annotated with publisher credibility metadata, user engagement graphs and temporal dissemination cascades.
- **Advantages:** Contains 23,196 fact-checked articles with authoritative ground-truth labels. Supports joint exploration of news text and social propagation. Establishes an industry-standard benchmark for social context-aware research.
- **Limitations:** Suffers from incomplete historical social media data due to API rate limits

and deleted posts. Heavy reliance on external social networks makes the dataset ineffective for evaluating newly published standalone articles.

3.4 Impact of Convolutional Neural Network and FastText Embedding on Text Classification

Authors: Muhammad Umer, Zainab Imtiaz, Muhammad Ahmad et al.

Published On: Multimedia Tools and Applications, 2023

- **About:** Investigates the combination of FastText subword embeddings and 1D Convolutional Neural Networks for classifying noisy and unstructured text.
- **Methodology:** Raw text is processed through FastText to capture character-level n-gram patterns. Parallel convolutional filters of varying kernel sizes extract local syntactic features, followed by global max-pooling and dropout regularisation.
- **Advantages:** Handles unknown vocabulary, typos and misspellings effectively. Significantly faster training and inference than Transformer models. Exceeds 94.5 % accuracy on benchmark classification tasks.
- **Limitations:** Focuses on localized phrase patterns and misses long-range semantic dependencies across long documents. Performance is sensitive to convolutional hyperparameter tuning.

3.5 FakeBERT: Fake News Detection in Social Media using BERT and Deep Learning

Authors: Rohit Kumar Kaliyar, Anurag Goswami and Pratik Narang

Published On: Soft Computing, Springer, 2021

- **About:** Proposes a deep-learning architecture combining bidirectional contextual representations from BERT with parallel convolutional layers for fake news detection.
- **Methodology:** Pre-trained BERT is fine-tuned to capture bidirectional semantics. The resulting embeddings pass through parallel 1D-CNN channels with different filter sizes before Softmax classification.
- **Advantages:** Models bidirectional sentence context alongside local semantic phrases. Achieves 98.90 % accuracy on benchmark corpora and is robust against subtle deceptive writing styles.

- **Limitations:** Requires heavy GPU compute and large memory bandwidth. Operates as an opaque black box with no decision transparency. High inference latency makes it unsuitable for sub-second web deployments.

3.6 A Systematic Survey on Explainable AI Applied to Fake News Detection

Authors: S. D. Madhu Kumar and Angel Mary Chacko

Published On: Engineering Applications of Artificial Intelligence, 2023

- **About:** Provides a structured analysis of Explainable AI frameworks applied to fake news detection, examining local, global and model-agnostic explanation approaches.
- **Methodology:** Categorises and evaluates post-hoc explanation techniques including LIME, SHAP, attention heatmaps and layer-wise relevance propagation across deep-learning architectures, assessing fidelity, stability and user trust.
- **Advantages:** Identifies key criteria for evaluating explanation quality. Highlights systemic biases in black-box fact-checking and emphasises the role of interpretable word attribution in building end-user credibility.
- **Limitations:** Perturbation-based explainers such as LIME introduce noticeable computational latency during real-time serving. No standardized quantitative metrics exist for benchmarking human trust in AI explanations.

3.7 Exploiting Stance Similarity and Graph Neural Networks for Fake News Detection

Authors: Kazuki Soga, Suguru Yoshida and Masahiro Muneyasu

Published On: Pattern Recognition Letters, Elsevier, 2024

- **About:** Investigates Graph Neural Networks that exploit user stance similarity and social propagation topology for news authenticity classification.
- **Methodology:** Constructs heterogeneous bipartite graphs capturing interactions between news articles and social media users. Edge weights reflect stance similarity from comments and Graph Convolutional Networks learn representation vectors over the graph structure.
- **Advantages:** Integrates collective crowd skepticism with news content. Improves F1-scores over standalone text models and shows enhanced resistance to adversarial text

perturbations.

- **Limitations:** Requires extensive user-interaction and social-engagement records. Needs substantial memory and compute for graph propagation. Cannot evaluate cold-start articles that have zero social commentary.

3.8 Big Data ML-Based Fake News Detection Using Distributed Learning

Authors: Abdullah Altheneyan and Abdullah Alhadlaq

Published On: IEEE Access, 2023

- **About:** Evaluates distributed machine-learning pipelines built on Apache Spark for scalable fake news classification over large news streams.
- **Methodology:** Uses Spark MLlib for parallelised TF-IDF feature extraction and N-gram generation. Trains and ensembles Random Forest, Naive Bayes and Gradient Boosted Trees classifiers across distributed clusters.
- **Advantages:** Processes large volumes of enterprise news data efficiently. Achieves 92.4 % accuracy on the FNC-1 benchmark and provides high throughput suitable for big-data pipelines.
- **Limitations:** TF-IDF and N-gram features miss deeper semantic relationships and sequential context. Setting up and maintaining distributed Spark clusters incurs high operational overhead. Predictions lack explanatory attribution.

3.9 “Liar, Liar Pants on Fire”: A New Benchmark Dataset for Fake News Detection

Authors: William Yang Wang

Published On: ACL, Association for Computational Linguistics, 2017

- **About:** Introduces the LIAR benchmark dataset providing fine-grained, multi-tier truthfulness ratings for political claim verification.
- **Methodology:** Compiles 12,800 short political statements from PolitiFact classified into six truthfulness levels. Evaluates standard machine-learning models, CNNs and bi-directional LSTMs using statement text and speaker metadata.

- **Advantages:** Provides an authoritative, human-annotated benchmark supporting nuanced multi-class veracity classification. Evaluates the impact of speaker history, party affiliation and context alongside lexical features.
- **Limitations:** Short statements frequently lack sufficient textual context for deep semantic understanding. Text-only models struggle without speaker metadata. Automated visual explanations are not provided.

3.10 Fake News Detection: A Hybrid CNN-RNN Based Approach

Authors: Jamal Abdul Nasir, Osama Subhani Khan and Iraklis Varlamis

Published On: International Journal of Information Management Data Insights, 2021

- **About:** Proposes a hybrid neural architecture combining Convolutional Neural Networks with Recurrent Neural Networks for automated news text classification.
- **Methodology:** Word embeddings pass through 1D-CNN layers to identify salient n-gram representations, followed by LSTM layers that preserve long-range contextual relationships. Dense layers generate binary veracity predictions.
- **Advantages:** Successfully models both localized phrase patterns and long-term sentence flow. Outperforms isolated CNN and RNN baselines on real-world corpora such as ISOT without requiring external social-media metadata.
- **Limitations:** Functions as a black-box model lacking built-in explanatory capabilities. Standard word embeddings struggle with vocabulary gaps and typos. Training latency increases on very long articles.

3.11 Supervised Learning for Fake News Detection: An Explainable Approach

Authors: Julio C. S. Reis, André Correia, Fabrício Murai, Adriano Veloso and Fabricio Benevenuto

Published On: IEEE Internet Computing, 2019

- **About:** Explores explainable supervised learning for fake news detection by quantifying the predictive importance of linguistic, emotional and syntactic features.
- **Methodology:** Extracts handcrafted feature groups covering lexical diversity, readability scores, sentiment polarity and grammatical structures. Machine-learning classifiers are

trained and global feature importance rankings are derived.

- **Advantages:** Identifies writing characteristics that strongly indicate deceptive text, confirming that sensational and emotionally charged vocabulary is a key marker of fake news. Provides globally interpretable feature rankings.
- **Limitations:** Relies heavily on manual feature engineering rather than automatic representation learning. Achieves lower accuracy on complex semantic sentences compared to modern neural networks. Cannot highlight suspicious words in individual articles in real time.

Chapter 4

SYSTEM DESIGN

System design serves as the blueprint for translating functional requirements and theoretical concepts into an operational software architecture. This chapter presents the architectural and design diagrams for the Ledger fake news detection system. The visual representations delineate the multi-tier system topology, the hierarchical movement of data through progressive processing abstractions and the chronological interaction lifecycle among distributed components.

4.1 System Architecture

The system architecture follows a modular three-tier structure comprising the presentation tier, the backend application tier and the core artificial intelligence engine tier. The architectural design guarantees loose coupling, distinct separation of concerns and efficient horizontal scalability.

The presentation tier is implemented as a responsive web dashboard using React.js. It provides an intuitive interface for journalists, fact-checkers and everyday users to submit arbitrary news articles or headlines for verification. The interface receives user input, performs client-side validation and renders the computed authenticity verdict alongside confidence ratings and an interactive word-importance heatmap.

The backend application tier is constructed with Python FastAPI and served via an asynchronous ASGI server. It exposes high-performance RESTful API endpoints for news verification requests. The backend validates incoming payloads, orchestrates data transfer between the frontend and the machine learning pipeline and formats output predictions into standardized JSON responses.

The core AI engine tier hosts the natural language processing preprocessor, the FastText embedding module, the hybrid 1D-CNN and LSTM neural network and the LIME explainability module. FastText subword embeddings represent input tokens as 300-dimensional dense vectors. The 1D-CNN layers extract localized n-gram syntax and deceptive phrase patterns, while the LSTM layers capture long-range contextual dependencies across

the article. Finally, the LIME explainer perturbs the input text and computes token-level attribution weights to explain the classification verdict.




Figure 4.1: System Architecture Diagram

4.2 Data Flow Diagram Level 0

The Data Flow Diagram Level 0, also designated as the context diagram, defines the system boundary by illustrating the primary interactions between external entities and the Ledger system as a single unified process.

Two principal external entities interact with the system: the User (or Fact-Checker) and the Machine Learning Engineer (or System Administrator). The user submits unverified news text or article headlines to the Ledger system. In return, the system provides an authenticity assessment comprising the predicted label (Real or Fake), an associated confidence percentage and word-level explainability highlights indicating the linguistic basis for the prediction.

The machine learning engineer supplies annotated benchmark datasets, such as WELFake, FakeNewsNet and LIAR, for offline model training and fine-tuning. The system

outputs training convergence metrics, cross-validation performance summaries and loss statistics to allow the engineer to monitor model reliability.




Figure 4.2: Data Flow Diagram: Level 0 (Context Diagram)

4.3 Data Flow Diagram Level 1

The Data Flow Diagram Level 1 decomposes the unified Ledger system into four modular sub-processes, establishing the precise transformation of data from raw user input to an explainable prediction.

Process 1.0 (Preprocess Text) receives raw text from the user and applies cleaning, tokenization and lemmatization routines. Noise such as extraneous whitespace, HTML tags and non-alphanumeric characters is removed to generate normalized token streams.

Process 2.0 (Generate Embeddings) takes the normalized tokens and retrieves corresponding subword vector representations from the pre-trained FastText embedding repository. This step ensures that out-of-vocabulary terms, typographical errors and internet slang are projected into meaningful semantic vector spaces.

Process 3.0 (Classify News) feeds the dense embedding tensors into the trained CNN-LSTM neural network. Pre-trained weights stored in the model repository evaluate local syntactic patterns and sequential semantic flow, producing class probabilities for the Real and Fake categories.

Process 4.0 (Generate Explanation) receives the prediction from the classification process and invokes the LIME explainer. By perturbing words within the input sample and fitting a localized linear surrogate model, the process calculates positive and negative attribution weights for individual tokens. The final package, combining the veracity verdict, confidence score and word-level attribution highlights, is transmitted back to the user interface.




Figure 4.3: Data Flow Diagram: Level 1

4.4 UML Sequence Diagram

The UML Sequence Diagram illustrates the chronological message exchange between system participants during a complete news verification lifecycle. The diagram models six collaborating lifelines: the User, the React.js Frontend, the FastAPI Backend, the Text Preprocessor and FastText Vectorizer, the CNN-LSTM Classifier and the LIME Explainer.

The sequence begins when the user submits news text and triggers the verification action on the web interface. The React frontend sends an asynchronous HTTP POST request to the FastAPI verification endpoint. The backend invokes the preprocessing and FastText vectorizer module, which tokenizes the text and generates dense subword tensors before returning them to the API handler.

The backend forwards the tensor representations to the CNN-LSTM classification model, which computes the class likelihoods and returns the binary verdict alongside the confidence percentage. To ensure transparency, the backend dispatches the input text and model handle to the LIME explainer. LIME performs iterative word masking, queries the model for surrogate inference on perturbed variants and calculates localized word importance weights.

Once attribution weights are computed, LIME returns the attribution dictionary to the backend. The FastAPI server packages the verdict, confidence score and word highlights into a structured JSON payload and transmits it to the frontend client. The React interface parses the payload and dynamically updates the dashboard, presenting the user with a clear authenticity verdict and an interactive highlighted textual explanation.




Figure 4.4: UML Sequence Diagram

REFERENCES

1. E. Hashmi, S. Y. Yayilgan, M. M. Yamin, S. Ali and M. Abomhara, “Advancing Fake News Detection: Hybrid Deep Learning With FastText and Explainable AI,” *IEEE Access*, vol. 12, pp. 44463–44481, 2024. DOI: 10.1109/ACCESS.2024.3381038.
2. P. K. Verma, P. Agrawal, I. Amorim and R. Prodan, “WELFake: Word Embedding Over Linguistic Features for Fake News Detection,” *IEEE Transactions on Computational Social Systems*, vol. 8, no. 4, pp. 881–893, Aug. 2021.
3. K. Shu, D. Mahudeswaran, S. Wang, D. Lee and H. Liu, “FakeNewsNet: A Data Repository with News Content, Social Context and Spatiotemporal Information for Studying Fake News on Social Media,” *Big Data*, vol. 8, no. 3, pp. 171–188, Jun. 2020.
4. M. Umer, Z. Imtiaz, M. Ahmad, M. Nappi, C. Medaglia, G. S. Choi and A. Mehmood, “Impact of Convolutional Neural Network and FastText Embedding on Text Classification,” *Multimedia Tools and Applications*, vol. 82, no. 4, pp. 5569–5585, Feb. 2023.
5. R. K. Kaliyar, A. Goswami and P. Narang, “FakeBERT: Fake News Detection in Social Media using BERT and Deep Learning,” *Soft Computing*, vol. 25, no. 18, pp. 11771–11788, Sep. 2021.
6. S. D. M. Kumar and A. M. Chacko, “A Systematic Survey on Explainable AI Applied to Fake News Detection,” *Engineering Applications of Artificial Intelligence*, vol. 122, Art. no. 106087, Jun. 2023.
7. K. Soga, S. Yoshida and M. Muneyasu, “Exploiting Stance Similarity and Graph Neural Networks for Fake News Detection,” *Pattern Recognition Letters*, vol. 177, pp. 26–32, Jan. 2024.
8. A. Altheneyan and A. Alhadlaq, “Big Data ML-Based Fake News Detection Using Distributed Learning,” *IEEE Access*, vol. 11, pp. 29447–29463, 2023.
9. W. Y. Wang, “‘Liar, Liar Pants on Fire’: A New Benchmark Dataset for Fake News Detection,” in *Proc. 55th Annu. Meeting Assoc. Comput. Linguistics (ACL)*, 2017, pp. 422–426.

10. J. A. Nasir, O. S. Khan and I. Varlamis, "Fake news detection: A hybrid CNN-RNN based approach," *Int. J. Inf. Manage. Data Insights*, vol. 1, no. 2, Art. no. 100032, Nov. 2021.
11. J. C. S. Reis, A. Correia, F. Murai, A. Veloso and F. Benevenuto, "Supervised Learning for Fake News Detection: An Explainable Approach," *IEEE Internet Computing*, vol. 23, no. 2, pp. 46–53, Mar./Apr. 2019.